

Safety & Performance in Generated Coordination Code

Christopher Esterhuyse



Contents

- Background
 - Coordination
 - Reo
 - Reo Compiler
 - Runtime Coordination
- Contribution Overview
- Code Generation
 - Rust Language
 - Translation Pipeline
 - Example
- Generated Object Behavior
 - Execution
 - Performance
 - Safety
- Protocol Adherence
 - Problem Explained
 - Static Checking
 - Example

Contents

- Background

- Coordination
- Reo
- Reo Compiler
- Runtime Coordination

- Contribution Overview

- Code Generation

- Rust Language
- Translation Pipeline
- Example

- Generated Object Behavior

- Execution
- Performance
- Safety

- Protocol Adherence

- Problem Explained
- Static Checking
- Example

Background ●○○ Coordination

- In concurrent programming: Coordination involves **interactions** between actors.

Background ●○○ Coordination

- In concurrent programming: Coordination involves **interactions** between actors.
- **Imperative** languages express sequences of actions, mutating the program state.
- Concurrent programs must mix code for **computation** and **coordination**. Results in coupling.

```
void work_x(int* msg) {  
    for(int i = 0;; i++) {  
        sem_wait(&b);  
        *msg = i;  
        sem_post(&a);  
    }  
}
```

```
void work_y(int* msg) {  
    for(;;) {  
        sem_wait(&a);  
        printf("%d\n", *msg);  
        sem_post(&b);  
    }  
}
```

Background ●○○ Reo

Coordination language: Express coordination in a system more abstractly.

- Reo language is **graphical**: defines a *connector* with nodes and relations (\approx hyperedges).
- Relations constrain how nodes interact, defining how data can **flow**.

Eg: primitive *sync* in graphical syntax:




Background ●○○ Reo

Coordination language: Express coordination in a system more abstractly.

- Reo language is **graphical**: defines a *connector* with nodes and relations (\approx hyperedges).
- Relations constrain how nodes interact, defining how data can **flow**.
- Nodes exposed to the environment are **ports**.
- Non-primitive connectors are **built** from others by exposing or coupling ports.


Eg: *alternator2* in graphical & textual syntax:



```
alternator2(A, B, C){  
  syncdrain(A, B)  
  fifo1(A, C)  
  sync(B, C)  
}
```


Background ●●○ Reo Compiler

- Tool for translating a Reo specification to a *protocol object* in a target language.



Background ●●○ Reo Compiler

- Tool for translating a Reo specification to a *protocol object* in a target language.



- Has *backends* for targeting Java, Maude, Promela, ...
 - The role of the protocol object depends on the language.
 - For Java: generate coordinating Java glue-code.

Background ●●○ Reo Compiler

- Tool for translating a Reo specification to a *protocol object* in a target language.




- Has *backends* for targeting Java, Maude, Promela, ...
 - The role of the protocol object depends on the language.
 - For Java: generate coordinating Java glue-code.

```
void work_x(int* msg) {  
    for(int i = 0;; i++) {  
        sem_wait(&b);  
        *msg = i;  
        sem_post(&a);  
    }  
}
```

Background ●●○ Reo Compiler

- Tool for translating a Reo specification to a *protocol object* in a target language.




- Has *backends* for targeting Java, Maude, Promela, ...
 - The role of the protocol object depends on the language.
 - For Java: generate coordinating Java glue-code.
- Developed at CWI. Has seen work by...
 - Sung-Shik Jongmans
 - Kasper Dokter
 - Benjamin Lion
 - ...

Background ●●○ Reo Compiler

Building a coordinating program with the Reo compiler:

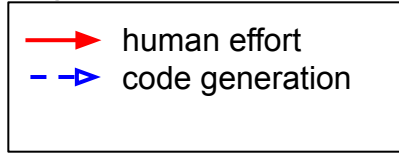
Background ●●○ Reo Compiler

Building a coordinating program with the Reo compiler:




Step 1: Express coordination logic in Reo.
Compile to protocol component(s).

legend:



Background ●●○ Reo Compiler

Building a coordinating program with the Reo compiler:




Step 2: Express computation logic as components.
Data exchange is abstracted to port operations with the environment.

```
void work_x(port p) {  
    for(int i = 0;; i++) {  
        p.put(i);  
    }  
}
```


Background ●●○ Reo Compiler

Building a coordinating program with the Reo compiler:




Step 3: Build program by linking components' ports as specified.

Eg: protocol coordinating 3 other components



Background ●●○ Reo Compiler

Building a coordinating program with the Reo compiler:




Advantages:

- **Direct:** Modify and check protocol via specification.
- **Indirect:** Program components are modular, reusable, maintainable.


Background ●●● Runtime Coordination

- User components initiate **put** or **get** actions on their local ports as they like.



Background ●●● Runtime Coordination

- User components initiate **put** or **get** actions on their local ports as they like.
- *Protocol object* completes sets of actions as soon as they can be organized into a specified *interaction*.




Contents

- Background
 - Coordination
 - Reo
 - Reo Compiler
 - Runtime Coordination
- Contribution Overview
- Code Generation
 - Rust Language
 - Translation Pipeline
 - Example
- Generated Object Behavior
 - Execution
 - Performance
 - Safety
- Protocol Adherence
 - Problem Explained
 - Static Checking
 - Example

Contribution Overview

Task: Make a new Reo compiler backend for a systems-language target.

Holy grail: Reo-generated code is as safe and performant as hand-crafted code.



Make it safe	Make it practical
Correct wrt. specification : protocol and general Reo semantics preserved.	Implement optimizations (eg. reference-passing port operations in shared memory).
Correct wrt. target language : code type-checks, no undefined behavior, etc.	Make it possible for programmers to reason about liveness .

Contents

- Background
 - Coordination
 - Reo
 - Reo Compiler
 - Runtime Coordination
- Contribution Overview
- Code Generation
 - Rust Language
 - Translation Pipeline
 - Example
- Generated Object Behavior
 - Execution
 - Performance
 - Safety
- Protocol Adherence
 - Problem Explained
 - Static Checking
 - Example

Code Generation ●○○ Rust language

We compile to Rust.

- Systems language
 - Low-level resource manipulations → optimizations.
 - Cheap interoperability with C++ and C → larger audience.



Code Generation ●○○ Rust language



We compile to Rust.

- Systems language
 - Low-level resource manipulations → optimizations.
 - Cheap interoperability with C++ and C → larger audience.

- Expressive affine type system
 - Expressive component and port API to statically enforce safety.

```
fn foo (x: & mut Bar)
```

requires exclusive access to argument. may mutate

x passed by reference

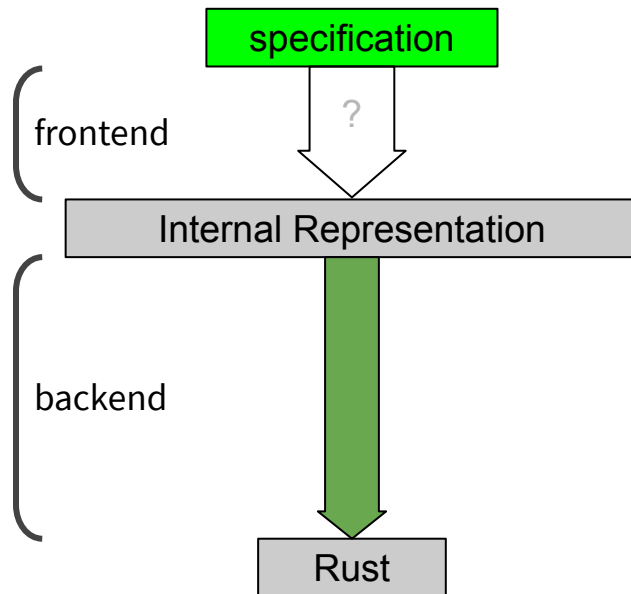
Code Generation ●○○ Rust language



We compile to Rust.

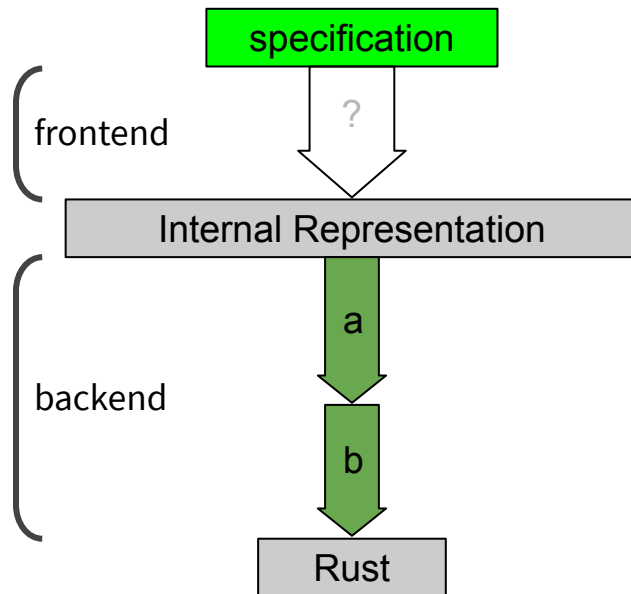
- Systems language
 - Low-level resource manipulations → optimizations.
 - Cheap interoperability with C++ and C → larger audience.
- Expressive affine type system
 - Expressive component and port API to statically enforce safety.
- Popular modern language
 - Lively, growing community → likely to stay a useful Reo target.
 - Modern features for safety and productivity (eg. closures, matching).

Code Generation ●●○ Translation Pipeline




Code Generation ●●○ Translation Pipeline

- Back-end translation can be broken into stages:
 - a. **Abstract:** Port data types resolved, interactions laid out as sequences of actions.
 - b. **Concrete:** Emit executable, concrete Rust.



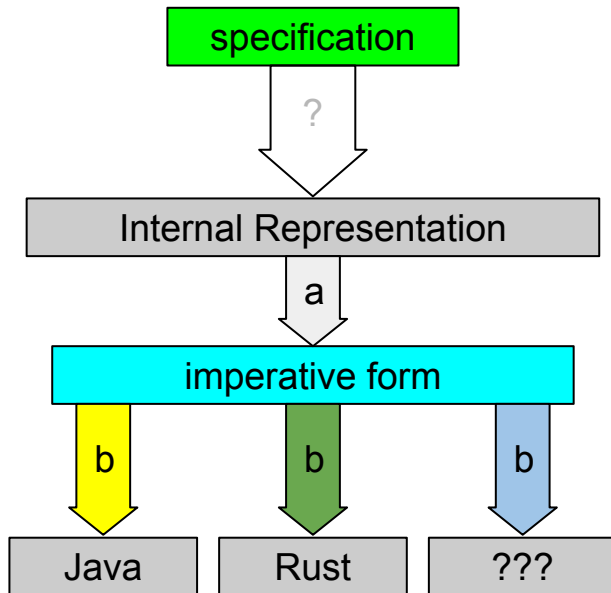
Code Generation ●●○ Translation Pipeline

- Back-end translation can be broken into stages:
 - a. **Abstract:** Port data types resolved, interactions laid out as sequences of actions.
 - b. **Concrete:** Emit executable, concrete Rust.
- Similar for all imperative languages.



Code Generation ●●○ Translation Pipeline

- Back-end translation can be broken into stages:
 - Abstract:** Port data types resolved, interactions laid out as sequences of actions.
 - Concrete:** Emit executable, concrete Rust.
- Similar for all imperative languages.
- Introduce *imperative form* between **a** and **b**.
 - Sensible for many imperative targets (eg. not promela).
 - Reo compiler does **a**. Rust does **b**.
→ Reo and Rust decoupled.



Note: only Rust currently does this.

Code Generation ●●● Example

User sees

```
pub fn protocol_1<X: Eq>() -> ProtoHandle {
```

Explicit data types

```
    let xtype = TypeInfo::of::<X>();
    let booltype = TypeInfo::of::<bool>();
    ProtoDef {
        name_defs: hashmap! {
            "A" => Port { is_putter:true, type_info: xtype },
            "B" => Port { is_putter:true, type_info: xtype },
            "m" => Mem ( booltype ),
        },
```

Interaction as action sequence


```
        rules: vec![RuleDef {
            state_guard: StatePredicate {
                ready_ports: hashset! { "A", "B" },
                full_mem: hashset! {},
                empty_mem: hashset! { "m" },
            },
            ins: vec![
                Check(Eq(Named("A"), Named("B")),
                CreateFromFormula { dest: "temp", term: True },
            ],
            output: hashmap! { "temp" => "m" },
        }],
    }.build(MemInitial::default()).unwrap()
}
```

Contents

- Background
 - Coordination
 - Reo
 - Reo Compiler
 - Runtime Coordination
- Contribution Overview
- Code Generation
 - Rust Language
 - Translation Pipeline
 - Example
- Generated Object Behavior
 - Execution
 - Performance
 - Safety
- Protocol Adherence
 - Problem Explained
 - Static Checking
 - Example

Generated Object Behavior ●○○ Execution


- Reo→Java:
 - Distinct **class**. Implements `Protocol` interface.
 - Protocol is threaded.
 - Data exchanged at ports.
 - Value's reference is moved.



Generated Object Behavior ●○○ Execution


- Reo→Java:

- Distinct **class**. Implements `Protocol` interface.
- Protocol is threaded.
- Data exchanged at ports.
- Value's reference is moved.



- Reo→Rust:


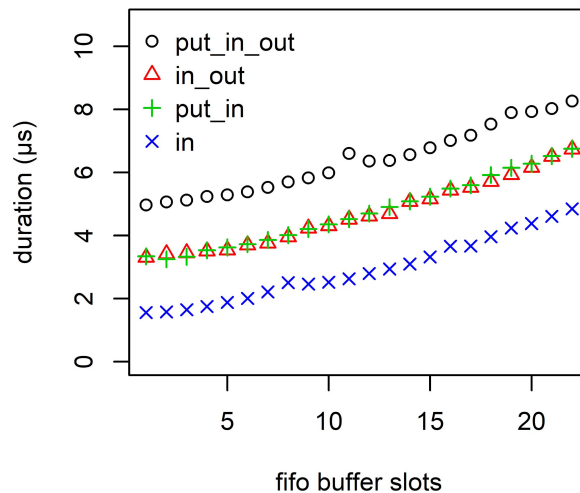
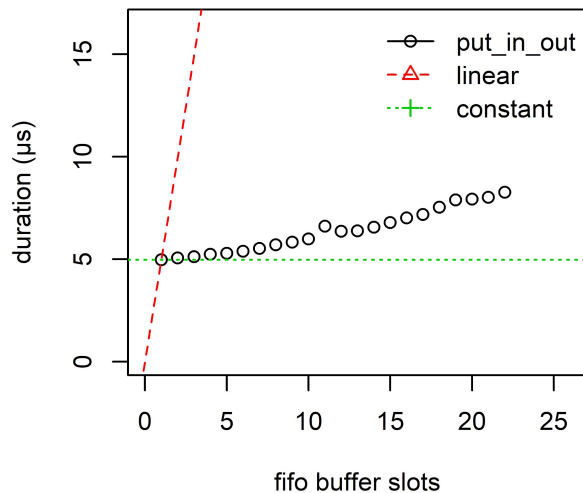
- Distinct **function**. Builds `Protocol` instance.
- Protocol is data. Boundary components do work.
- Data acquired from the source (eg. $C_0 \rightarrow C_1$).
- Value is moved.



Generated Object Behavior ●●○ Performance


- Optimization 1: internal reference-passing and aliasing
 - Values moved *within* protocol object by reference.
 - Cost of internal move or replication is a small constant.

fifo-N connector:




Generated Object Behavior ●●○ Performance

- Naïvely, interactions occur sequentially
 - Interactions influence protocol's state.
→ they must be serializable




Generated Object Behavior ●●○ Performance

- Naïvely, interactions occur sequentially
 - Interactions influence protocol's state.
→ they must be serializable
- Optimization 2 : concurrent data movements
 - Interactions are still serialized, but their work may overlap



Generated Object Behavior ●●○ Performance

- Naïvely, interactions occur sequentially
 - Interactions influence protocol's state.
→ they must be serializable
- Optimization 2 : concurrent data movements
 - Interactions are still serialized, but their work may overlap

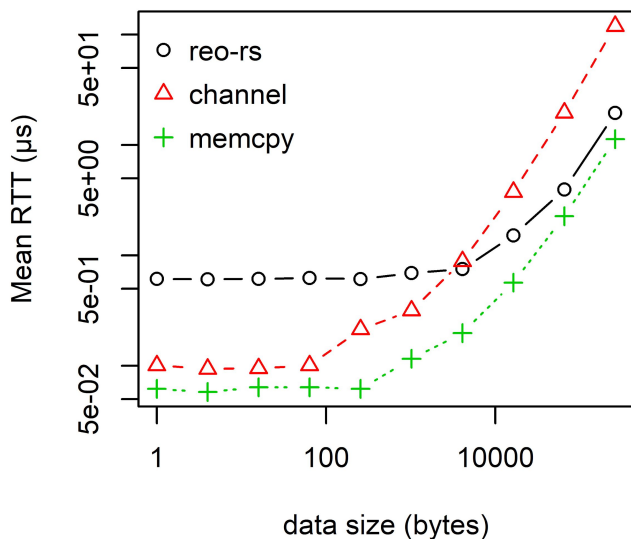


<u>3 putter-getter pairs:</u>	mean active time			run duration	mean parallelism
	p0	p1	p2		
move	2.68μs	2.594μs	2.993μs	31.1705ms	2.652
copy	2.737μs	2.4μs	2.673μs	28.7161ms	2.720
signal	2.351μs	2.282μs	1.943μs	24.7852ms	2.653
clone	4.451ms	4.461ms	4.416ms	44.609s	2.988

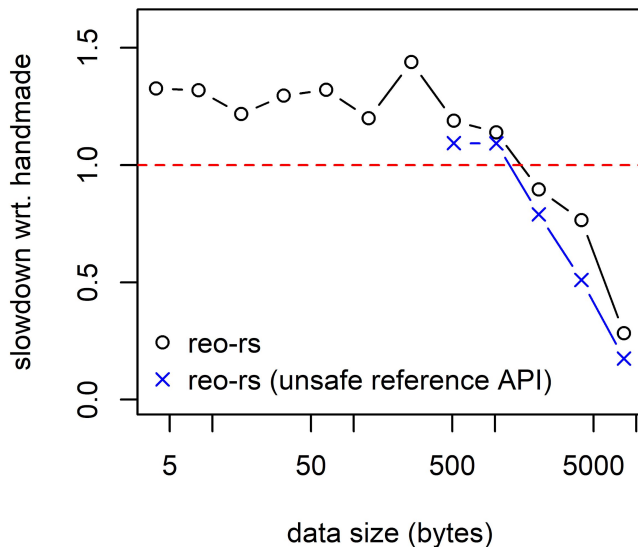
Generated Object Behavior ●●○ Performance

Competitive vs. hand-crafted code for non-trivial protocols.

1-thread fifo1:



alternator2:



Generated Object Behavior ●●● Safety

Port & protocol objects are always in valid state.

- Protocol objects only acquired with `build`, which ensures...
 - Initialization is beyond user's control (ie. port linkage and initial state)

```
fn build(&ProtoDef, MemInitial)  
-> Result<ProtoHandle, BuildError>
```

returned object is initialized

Generated Object Behavior ●●● Safety

Port & protocol objects are always in valid state.

- Protocol objects only acquired with `build`, which ensures...
 - Initialization is beyond user's control (ie. port linkage and initial state)
- Ports acquired only acquired from a protocol with `claim`, which ensures...
 - port objects are of the correct type (data type, orientation).
 - cannot duplicate logical ports.

```
fn claim<T>(& ProtoHandle, Name)  
-> Result<Putter<T>, ClaimError>
```

may return error

port data type and orientation is fixed

Generated Object Behavior ●●● Safety

Reo's semantics are preserved.

- Port **put** and **get** use Rust's ownership semantics to ensure...
 - values are passed as expected.
 - a port cannot do two things at once.

```
fn get (& mut Port) -> T
```

operation requires exclusive access

Independent value returned

Generated Object Behavior ●●● Safety

Reo's semantics are preserved.

- Port **put** and **get** use Rust's ownership semantics to ensure...
 - values are passed as expected.
 - a port cannot do two things at once.
- Replication via user-defined **clone** operation.
 - Signature prohibits side effects → more optimizations possible.

```
fn clone (& T) -> T
```

original is immutable

independent value returned

Contents


- Background
 - Coordination
 - Reo
 - Reo Compiler
 - Runtime Coordination
- Contribution Overview
- Code Generation
 - Rust Language
 - Translation Pipeline
 - Example
- Generated Object Behavior
 - Execution
 - Performance
 - Safety
- Protocol Adherence
 - Problem Explained
 - Static Checking
 - Example

Protocol Adherence ●○○ Problem Explained

- The behavior of a system is the *composition* of the behavior of its components. System behavior is constrained by the protocol.

Protocol Adherence ●○○ Problem Explained

- The behavior of a system is the *composition* of the behavior of its components. System behavior is constrained by the protocol.



- Eg. this system is specified to have no behavior.


Protocol Adherence ●○○ Problem Explained

- Recall: Components are ignorant of their environment.



Protocol Adherence ●○○ Problem Explained


- Recall: Components are ignorant of their environment.



- Components define **local** behavior, but we observe **global** behavior at runtime.
 - We don't know the effect of local actions on the system's *liveness*.


Protocol Adherence ●●○ Static Checking

- **Governor:** Enforces *adherence*, i.e., local actions don't add new constraints.



Protocol Adherence ●●○ Static Checking

- **Governor:** Enforces *adherence*, i.e., local actions don't add new constraints.
 - constructed from the specification of the connected protocol (but has a localized view).



Protocol Adherence ●●○ Static Checking

- We focus on **static** governors. Deviation is checked at compile-time.

Protocol Adherence ●●○ Static Checking

- We focus on **static** governors. Deviation is checked at compile-time.
- Encode governors into Rust's type system.
 - Component is adherent \Leftrightarrow Rust code compiles.
 - **Why:** Programmer needs the Rust compiler anyway. Governor is only a Rust dependency.
 - **How:** Translate “port operation is permitted in current state” to “variable has correct type”

Protocol Adherence ●●○ Static Checking

- We focus on **static** governors. Deviation is checked at compile-time.
- Encode governors into Rust's type system.
 - Component is adherent \Leftrightarrow Rust code compiles.
 - **Why:** Programmer needs the Rust compiler anyway. Governor is only a Rust dependency.
 - **How:** Translate “port operation is permitted in current state” to “variable has correct type”:
 - component always has one *token*. Its type encodes the protocol's state.
 - port operations replace the token \Leftrightarrow protocol's state is changed.

before: `fn get (& mut Port,) -> T`

after: `fn get (& mut Port, X) -> (T, Y)`

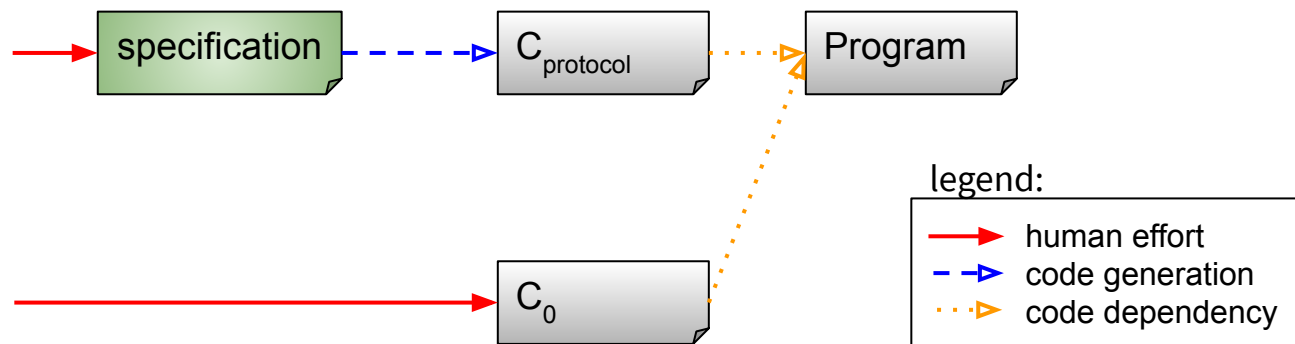
argument consumed

Initialized value returned

Protocol Adherence ●●○ Static Checking

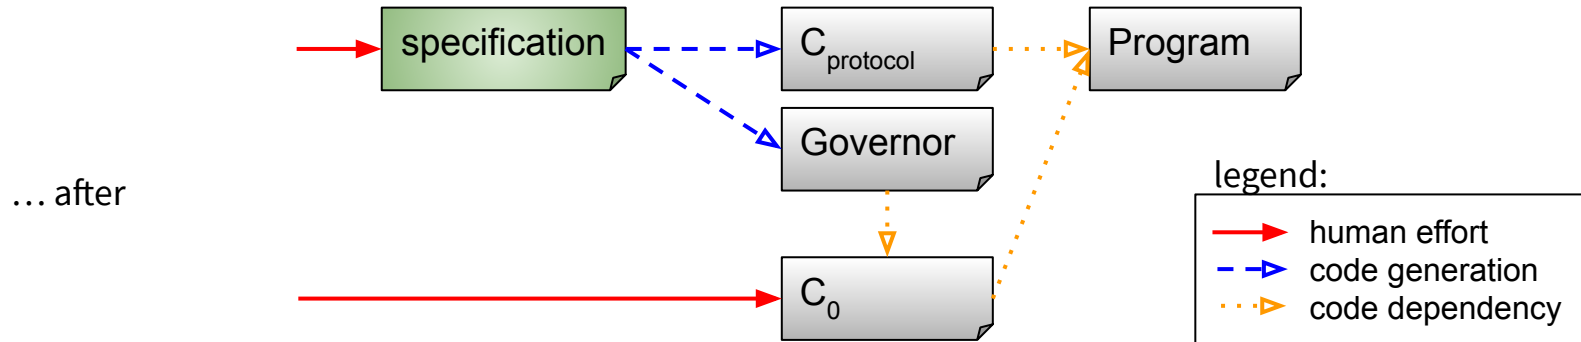
Governors supplement the user's workflow:

before...



Protocol Adherence ●●○ Static Checking

Governors supplement the user's workflow:



Protocol Adherence ●●● Example

```
fn unsafe_component_0(port_a, port_b) {  
    DATA = port_a.get();  
    port_b.put(DATA);  
    port_b.put(DATA); // (deadlocks at runtime)  
}
```

Protocol Adherence ●●● Example

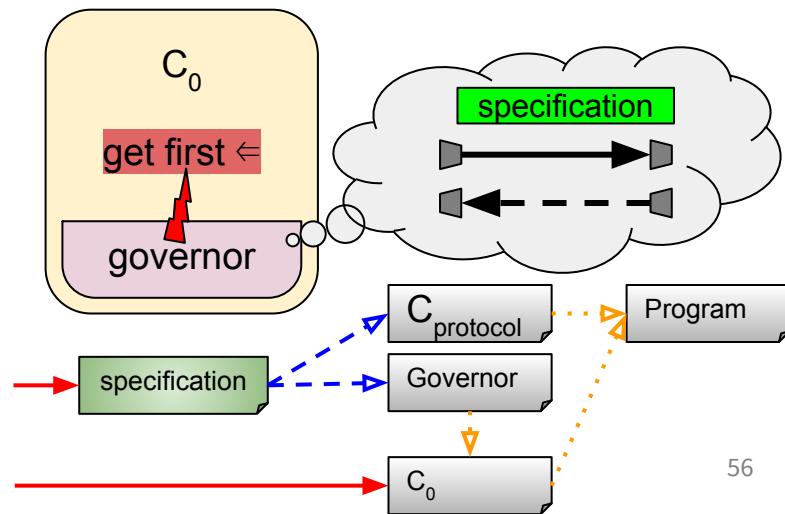
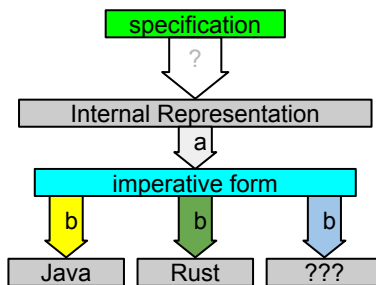



```
fn unsafe_component_0(port_a, port_b) {  
    DATA = port_a.get();  
    port_b.put(DATA);  
    port_b.put(DATA); // (deadlocks at runtime)  
}
```

```
fn __safe_component_0(t, port_a, port_b) {  
    (t, DATA) = port_a.get(t);  
    t = port_b.put(t, DATA);  
    t = port_b.put(t, DATA);  
    // TYPE ERROR  ^-- Expected FOO not BAR.  
}
```

End Slide: Summary

- Contributed:
 - Reo compiler **backend** to Rust.
 - Rust library for generic protocol types and behaviors (eg: **build**).
 - Performance **benchmarks** for Rust protocol objects.
 - Design of static **governors** in Rust's type system.

	mean active time			run duration	mean parallelism
	p0	p1	p2		
move	2.68μs	2.594μs	2.993μs	31.1705ms	2.652
copy	2.737μs	2.4μs	2.673μs	28.7161ms	2.720
signal	2.351μs	2.282μs	1.943μs	24.7852ms	2.653
clone	4.451ms	4.461ms	4.416ms	44.609s	2.988



Extra Slides

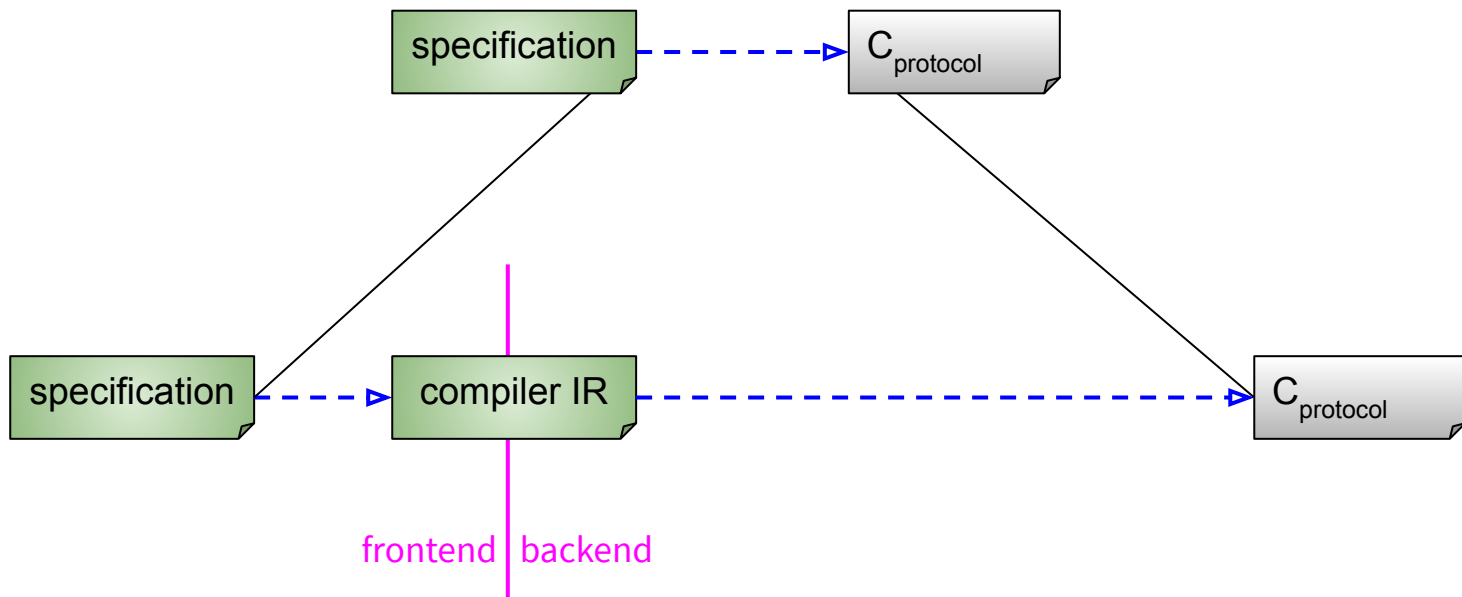
Extra ●○ Translation Pipeline

- Translation from Reo to Rust takes several steps



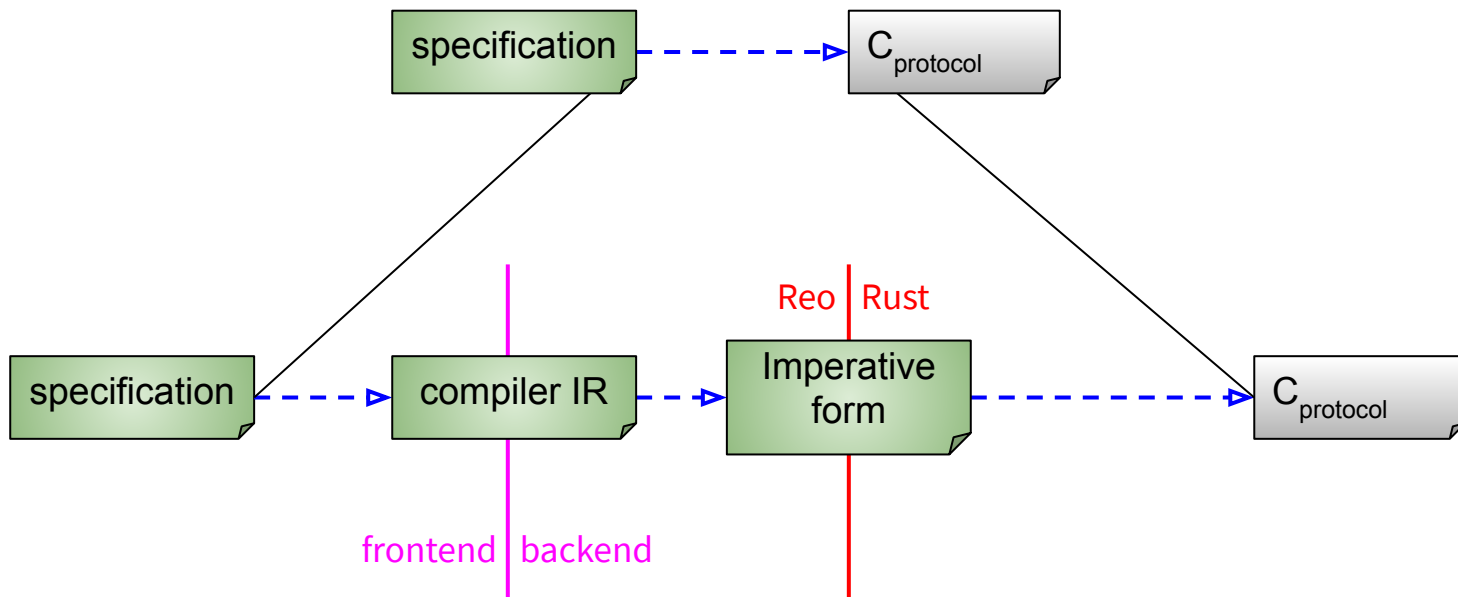
Extra ●○ Translation Pipeline

- Translation from Reo to Rust takes several steps
 - This work takes the **frontend**'s work as given



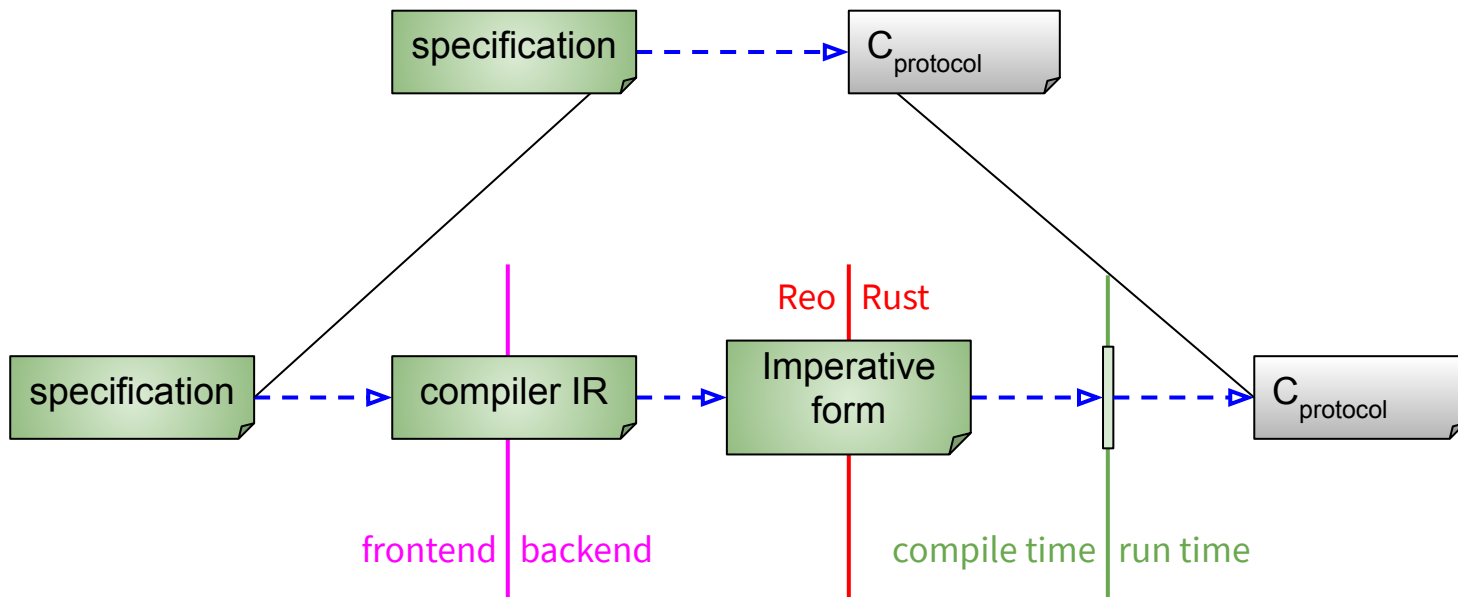
Extra ●○ Translation Pipeline

- Translation from Reo to Rust takes several steps
 - This work takes the **frontend**'s work as given
 - Imperative form represents translation *before* rust-specifics + optimization




Extra ●○ Translation Pipeline

- Translation from Reo to Rust takes several steps
 - This work takes the **frontend**'s work as given
 - Imperative form represents translation *before* rust-specifics + optimization
 - Allows us to finish translation at runtime (for example)



Extra ●● Imperative Form

- New intermediate form for protocols
 - Between Reo specification and generated target code
 - Interactions → sequences of abstract actions



translate

abort if $\text{ready} \subset \{A, B, C\}$

abort if M is full

move $\{A \rightarrow M, B \rightarrow C\}$